

# Multiprocessing

---



**Tim Ojo**

@tim\_ojo [www.timojo.com](http://www.timojo.com)



# Processes vs. Threads


---




# Process

a running instance of a computer program








Process 1



Process 2







# Processes vs. Threads

- **Sidesteps GIL**





# Processes vs. Threads

- **Sidesteps GIL**
- **Less need for synchronization**



# Processes vs. Threads

- Sidesteps GIL
- Less need for synchronization
- Can be paused and terminated



# Processes vs. Threads

- Sidesteps GIL
- Less need for synchronization
- Can be paused and terminated
- More resilient



# Processes vs. Threads

- **Sidesteps GIL**
- **Less need for synchronization**
- **Can be paused and terminated**
- **More resilient**

- **Higher memory footprint**
- **Expensive context switches**



# Multiprocessing API

---



```
1:  import threading

2:  def do_some_work(val):
3:      print ("doing some work in thread")
4:      print ("echo: {}".format(val))
5:      return

6:  val = "text"
7:  t = threading.Thread(target=do_some_work, args=(val,))
8:  t.start()
9:  t.join()
10:
```

```
1:  import multiprocessing

2:  def do_some_work(val):
3:      print ("doing some work in thread")
4:      print ("echo: {}".format(val))
5:      return

6:  if __name__ == '__main__':
7:      val = "text"
8:      t = multiprocessing.Process(target=do_some_work, args=(val,))
9:      t.start()
10:     t.join()
```

# Pickling

is the process whereby a Python object hierarchy is converted into a byte stream. “unpickling” is the inverse operation





# Picklable Objects

**None, True, False**

**Integers, floats, complex numbers**

**Normal and Unicode strings**

**Collections containing only picklable objects**

**Top level functions**

**Classes with picklable attributes**



```
1:  import multiprocessing

2:  def do_some_work(val):
3:      print ("doing some work in thread")
4:      print ("echo: {}".format(val))
5:      return

6:  if __name__ == '__main__':
7:      val = "text"
8:      t = multiprocessing.Process(target=do_some_work, args=(val,))
9:      t.start()
10:     t.join()
```

```
import multiprocessing

print("outer: this will print when run and when imported")

def func():
    print("func: this will print only when run")
    return

if __name__ == '__main__':
    print("main: this will print only when run")
    p = multiprocessing.Process(target=func)
    p.start()
    p.join()
```

```
class multiprocessing.Process(group=None,  
                             target=None,  
                             name=None,  
                             args=(),  
                             kwargs={},  
                             daemon=None)
```



```
class multiprocessing.Process(group=None,  
                             target=None,  
                             name=None,  
                             args=(),  
                             kwargs={},  
                             daemon=None)
```



```
class multiprocessing.Process(group=None,  
                             target=None,  
                             name=None,  
                             args=(),  
                             kwargs={},  
                             daemon=None)
```



```
class multiprocessing.Process(group=None,  
                             target=None,  
                             name=None,  
                             args=(),  
                             kwargs={},  
                             daemon=None)
```



```
class multiprocessing.Process(group=None,  
                             target=None,  
                             name=None,  
                             args=(),  
                             kwargs={},  
                             daemon=None)
```





```
class multiprocessing.Process(group=None,  
                             target=None,  
                             name=None,  
                             args=(),  
                             kwargs={},  
                             daemon=None)
```



```
class multiprocessing.Process(group=None,  
                             target=None,  
                             name=None,  
                             args=(),  
                             kwargs={},  
                             daemon=None)
```



# Daemon Process

is a child process that does not prevent its parent process from exiting



# Terminating Processes



# Terminating Processes

- `is_alive()`
- `terminate()`



```
1: import multiprocessing
2: import time
3:
4: def do_work():
5:     print ('Starting do_work function')
6:     time.sleep(5)
7:     print ('Finished do_work function')
8:
9: if __name__ == '__main__':
10:     p = multiprocessing.Process(target=func)
11:     print ("[Before Start] Process is alive: {}".format(p.is_alive()))
12:     p.start()
13:     print ("[Running] Process is alive: {}".format(p.is_alive()))
14:     p.terminate()
15:     p.join()
16:     print ("[After Termination] Process is alive: {}".format(p.is_alive()))
17:     print ("Process exit code: {}".format(p.exitcode))
```



```
1: import multiprocessing
2: import time
3:
4: def do_work():
5:     print ('Starting do_work function')
6:     time.sleep(5)
7:     print ('Finished do_work function')
8:
9: if __name__ == '__main__':
10:     p = multiprocessing.Process(target=func)
11:     print ("[Before Start] Process is alive: {}".format(p.is_alive()))
12:     p.start()
13:     print ("[Running] Process is alive: {}".format(p.is_alive()))
14:     p.terminate()
15:     p.join()
16:     print ("[After Termination] Process is alive: {}".format(p.is_alive()))
17:     print ("Process exit code: {}".format(p.exitcode))
```



```
1: import multiprocessing
2: import time
3:
4: def do_work():
5:     print ('Starting do_work function')
6:     time.sleep(5)
7:     print ('Finished do_work function')
8:
9: if __name__ == '__main__':
10:     p = multiprocessing.Process(target=func)
11:     print ("[Before Start] Process is alive: {}".format(p.is_alive()))
12:     p.start()
13:     print ("[Running] Process is alive: {}".format(p.is_alive()))
14:     p.terminate()
15:     p.join()
16:     print ("[After Termination] Process is alive: {}".format(p.is_alive()))
17:     print ("Process exit code: {}".format(p.exitcode))
```





```
1: import multiprocessing
2: import time
3:
4: def do_work():
5:     print ('Starting do_work function')
6:     time.sleep(5)
7:     print ('Finished do_work function')
8:
9: if __name__ == '__main__':
10:     p = multiprocessing.Process(target=func)
11:     print ("[Before Start] Process is alive: {}".format(p.is_alive()))
12:     p.start()
13:     print ("[Running] Process is alive: {}".format(p.is_alive()))
14:     p.terminate()
15:     p.join()
16:     print ("[After Termination] Process is alive: {}".format(p.is_alive()))
17:     print ("Process exit code: {}".format(p.exitcode))
```



```
1: import multiprocessing
2: import time
3:
4: def do_work():
5:     print ('Starting do_work function')
6:     time.sleep(5)
7:     print ('Finished do_work function')
8:
9: if __name__ == '__main__':
10:     p = multiprocessing.Process(target=func)
11:     print ("[Before Start] Process is alive: {}".format(p.is_alive()))
12:     p.start()
13:     print ("[Running] Process is alive: {}".format(p.is_alive()))
14:     p.terminate()
15:     p.join()
16:     print ("[After Termination] Process is alive: {}".format(p.is_alive()))
17:     print ("Process exit code: {}".format(p.exitcode))
```



Process.terminate() gotchas



# Process.terminate() gotchas

- Shared resources may be put in an inconsistent state



# Process.terminate() gotchas

- Shared resources may be put in an inconsistent state
- Finally clauses and exit handlers will not be run



# Process Pools

---



```
class multiprocessing.Pool([num_processes  
                             [, initializer  
                             [, initargs  
                             [, maxtasksperchild ]]]] ))
```



```
class multiprocessing.Pool([num_processes  
                           [, initializer  
                           [, initargs  
                           [, maxtasksperchild ]]]) )
```





```
class multiprocessing.Pool([num_processes  
                             [, initializer  
                             [, initargs  
                             [, maxtasksperchild ]]]] ))
```



```
class multiprocessing.Pool([num_processes  
                             [, initializer  
                             [, initargs # picklable not required  
                             [, maxtasksperchild ]]]) )
```



```
class multiprocessing.Pool([num_processes  
                           [, initializer  
                           [, initargs  
                           [, maxtasksperchild ]]]] ))
```



```
map(func, iterable[, chunksize])
```



```
1: import multiprocessing
2:
3: def do_work(data):
4:     return data**2
5:
6: def start_process():
7:     print('Starting', multiprocessing.current_process().name)
8:
9: if __name__ == '__main__':
10:     pool_size = multiprocessing.cpu_count() * 2
11:     pool = multiprocessing.Pool(processes=pool_size,
12:                                initializer=start_process)
13:     inputs = list(range(10))
14:     outputs = pool.map(do_work, inputs)
15:     print('Outputs :', outputs)
16:
17:     pool.close() # no more tasks
18:     pool.join() # wait for the worker processes to exit
```



```
1: import multiprocessing
2:
3: def do_work(data):
4:     return data**2
5:
6: def start_process():
7:     print('Starting', multiprocessing.current_process().name)
8:
9: if __name__ == '__main__':
10:     pool_size = multiprocessing.cpu_count() * 2
11:     pool = multiprocessing.Pool(processes=pool_size,
12:                                initializer=start_process)
13:     inputs = list(range(10))
14:     outputs = pool.map(do_work, inputs)
15:     print('Outputs :', outputs)
16:
17:     pool.close() # no more tasks
18:     pool.join() # wait for the worker processes to exit
```



```
1: import multiprocessing
2:
3: def do_work(data):
4:     return data**2
5:
6: def start_process():
7:     print('Starting', multiprocessing.current_process().name)
8:
9: if __name__ == '__main__':
10:     pool_size = multiprocessing.cpu_count() * 2
11:     pool = multiprocessing.Pool(processes=pool_size,
12:                                initializer=start_process)
13:     inputs = list(range(10))
14:     outputs = pool.map(do_work, inputs)
15:     print('Outputs :', outputs)
16:
17:     pool.close() # no more tasks
18:     pool.join() # wait for the worker processes to exit
```



```
1: import multiprocessing
2:
3: def do_work(data):
4:     return data**2
5:
6: def start_process():
7:     print('Starting', multiprocessing.current_process().name)
8:
9: if __name__ == '__main__':
10:     pool_size = multiprocessing.cpu_count() * 2
11:     pool = multiprocessing.Pool(processes=pool_size,
12:                                initializer=start_process)
13:     inputs = list(range(10))
14:     outputs = pool.map(do_work, inputs)
15:     print('Outputs :', outputs)
16:
17:     pool.close() # no more tasks
18:     pool.join() # wait for the worker processes to exit
```





```
1: import multiprocessing
2:
3: def do_work(data):
4:     return data**2
5:
6: def start_process():
7:     print('Starting', multiprocessing.current_process().name)
8:
9: if __name__ == '__main__':
10:     pool_size = multiprocessing.cpu_count() * 2
11:     pool = multiprocessing.Pool(processes=pool_size,
12:                                initializer=start_process)
13:     inputs = list(range(10))
14:     outputs = pool.map(do_work, inputs)
15:     print('Outputs :', outputs)
16:
17:     pool.close() # no more tasks
18:     pool.join() # wait for the worker processes to exit
```



```
1: import multiprocessing
2:
3: def do_work(data):
4:     return data**2
5:
6: def start_process():
7:     print('Starting', multiprocessing.current_process().name)
8:
9: if __name__ == '__main__':
10:     pool_size = multiprocessing.cpu_count() * 2
11:     pool = multiprocessing.Pool(processes=pool_size,
12:                                initializer=start_process)
13:     inputs = list(range(10))
14:     outputs = pool.map(do_work, inputs)
15:     print('Outputs :', outputs)
16:
17:     pool.close() # no more tasks
18:     pool.join() # wait for the worker processes to exit
```



# Demo



## Process Pool Demo



`map_async(func, iterable[, chunksize[, callback]])` returns `AsyncResult`



```
map_async(func, iterable[, chunksize[, callback]]) returns AsyncResult
```

```
AsyncResult.get([timeout]) # returns the result when it arrives
```



```
apply(func[, args[, kwargs]])
```

```
apply_async(func[, args[, kwargs[, callback[, error_callback]]]])
```



```
1:  from multiprocessing import Pool
2:  import time
3:
4:  def multiply(x, y):
5:      return x*y
6:
7:  if __name__ == '__main__':
8:      pool = Pool(processes=4)
9:      result = pool.apply_async(mult, (7,7)) # evaluate "multiply(7,7)"
                                              asynchronously in a single process
10:     print result.get()  # prints 49
```



# Inter-process Communication

---






Inter-process  
Communication  
Channels


**Pipe**  
**Queue**



# multiprocessing.Pipe



# multiprocessing.Pipe



```
1: def make_tuple(conn):
2:     num = random.randint(1, 9)
3:     conn.send(('Hi ', num))
4:     print(conn.recv())
5:
6: def make_string(conn):
7:     tup = conn.recv()
8:     result = ""
9:     substr, num = tup
10:    for _ in range(num):
11:        result += substr
12:    conn.send(result)
13:
14: if __name__ == '__main__':
15:     conn1, conn2 = Pipe(True)
16:     p1 = Process(target=make_tuple, args=(conn1,))
17:     p2 = Process(target=make_string, args=(conn2,))
18:     p1.start()
19:     p2.start()
```





```
1: def make_tuple(conn):
2:     num = random.randint(1, 9)
3:     conn.send(('Hi ', num))
4:     print(conn.recv())
5:
6: def make_string(conn):
7:     tup = conn.recv()
8:     result = ""
9:     substr, num = tup
10:    for _ in range(num):
11:        result += substr
12:    conn.send(result)
13:
14: if __name__ == '__main__':
15:     conn1, conn2 = Pipe(True)
16:     p1 = Process(target=make_tuple, args=(conn1,))
17:     p2 = Process(target=make_string, args=(conn2,))
18:     p1.start()
19:     p2.start()
```



```
1: def make_tuple(conn):
2:     num = random.randint(1, 9)
3:     conn.send(('Hi ', num))
4:     print(conn.recv())
5:
6: def make_string(conn):
7:     tup = conn.recv()
8:     result = ''
9:     substr, num = tup
10:    for _ in range(num):
11:        result += substr
12:    conn.send(result)
13:
14: if __name__ == '__main__':
15:     conn1, conn2 = Pipe(True)
16:     p1 = Process(target=make_tuple, args=(conn1,))
17:     p2 = Process(target=make_string, args=(conn2,))
18:     p1.start()
19:     p2.start()
```



# Pipe vs. Queue



# Queue Methods

<b>threading .Queue</b>	<b>multiprocessing .Queue</b>
qsize()	qsize()
put()	put()
get()	get()
empty()	empty()
full()	full()
task_done()	<del>task_done()</del>
join()	<del>join()</del>





# Queue Methods

<b>threading .Queue</b>	<b>multiprocessing .Queue</b>	<b>multiprocessing .JoinableQueue</b>
qsize()	qsize()	qsize()
put()	put()	put()
get()	get()	get()
empty()	empty()	empty()
full()	full()	full()
task_done()	<del>task_done()</del>	task_done()
join()	<del>join()</del>	join()



```
1: def make_tuple(queue):
2:     num = random.randint(1, 9)
3:     queue.put(('Hi ', num))
4:     print(queue.get())
5:
6: def make_string(queue):
7:     tup = queue.get()
8:     result = ""
9:     substr, num = tup
10:    for _ in range(num):
11:        result += substr
12:    queue.put(result)
13:
14: if __name__ == '__main__':
15:     queue = Queue()
16:     p1 = Process(target=make_tuple, args=(queue,))
17:     p2 = Process(target=make_string, args=(queue,))
18:     p1.start()
19:     p2.start()
```



```
Queue.get([block[, timeout]])
```



```
Queue.put(obj,[[block[, timeout]])
```



# Demo



## Inter-Process Communication Demo



# Sharing State Between Processes

---



Shared State

# Shared Memory Manager Process



# Shared Memory

- multiprocessing.Value
- multiprocessing.Array





# Shared Memory

```
multiprocessing.Value(typecode_or_type, *args[, lock])
```



# Shared Memory

```
multiprocessing.Value(typecode_or_type, *args[, lock])
```



# ctypes

ctypes type	C type	Python type	Type code
c_bool	_Bool	bool(1)	
c_char	char	1-character bytes object	'c'
c_wchar	wchar_t	1-character string	'u'
c_int	int	int	'i'
c_long	long	int	'l'
c_float	float	float	'f'
c_char_p	char * (NUL terminated)	bytes object or None	
c_wchar_p	wchar_t * (NUL terminated)	string object or None	
c_void_p	void *	int or None	



```
counter = Value('i') # shared object of type int, defaults to 0

is_running = Value(ctypes.c_bool, False, lock=False) # shared
object of type boolean, defaulting to False, unsynchronized

my_lock = multiprocessing.Lock()
size_counter = Value('l', 0, lock=my_lock) # shared object of
type long, with a lock specified
```

```
counter = Value('i') # shared object of type int, defaults to 0
```

```
is_running = Value(ctypes.c_bool, False, lock=False) # shared  
object of type boolean, defaulting to False, unsynchronized
```

```
my_lock = multiprocessing.Lock()
```

```
size_counter = Value('l', 0, lock=my_lock) # shared object of  
type long, with a lock specified
```



```
counter = Value('i') # shared object of type int, defaults to 0
```

```
is_running = Value(ctypes.c_bool, False, lock=False) # shared  
object of type boolean, defaulting to False, unsynchronized
```

```
my_lock = multiprocessing.Lock()
```





```
size_counter = Value('l', 0, lock=my_lock) # shared object of  
type long, with a lock specified
```




# Manager




Value	Array
Dictionary	List
Lock	Semaphore




# Manager



Value	Array
Dictionary	List
Lock	Semaphore



Proxy




Proxy







# Manager



Value	Array
Dictionary	List
Lock	Semaphore



Proxy



Proxy



# Manager

```
multiprocessing.Manager() # spins up a new process
```



# Manager Shared Objects

Data Structures	Synchronization Mechanisms
Value	Lock
Array	RLock
List	BoundedSemaphore
Dict	Event
Namespace	Condition
Queue	Barrier



```
1: import multiprocessing
2:
3: def do_work(dictionary, item):
4:     dictionary[item] = item ** 2
5:
6: if __name__ == '__main__':
7:     mgr = multiprocessing.Manager()
8:     d = mgr.dict()
9:     jobs = [
10:         multiprocessing.Process(target=do_work, args=(d, i))
11:         for i in range(8)
12:     ]
13:     for j in jobs:
14:         j.start()
15:     for j in jobs:
16:         j.join()
17:     print('Results:', d)
```



# Demo



## Shared State Demo



# Process Synchronization

---



# Process Synchronization

```
with self.resized_size.get_lock():  
    self.resized_size += os.path.getsize(out_filepath)
```



# Process Synchronization

threading.	multiprocessing.
Lock	Lock
RLock	RLock
Semaphore	Semaphore
BoundedSemaphore	BoundedSemaphore
Event	Event
Condition	Condition
Barrier	Barrier





# Summary



Processes vs. Threads

Python Multiprocessing API

Process Pool

Inter-process Communication

Shared State Between Processes

Process Synchronization

